

HWpca_Fama_Perlo

January 13, 2025

1 Computational Linear Algebra: PCA Homework

1.1 Initialization:

Fill the missing values in this text box and in the following code-cell.

Academic Year: 2024/2025

1.1.1 Team Members (Alphabetical Order):

1. Famà, Daniele (345153);
2. Perlo, Alessandro (337131);

```
[1]: StudentID1 = 345153 # <----- Fill in the missing value
      StudentID2 = 337131 # <----- Fill in the missing value
```

1.2 Starting Code-Cell

1.2.1 Attention: DO NOT CHANGE THE CODE INSIDE THE FOLLOWING CELL!

```
[2]: #####
      ##### DO NOT CHANGE THE CODE IN THIS CELL #####
      #####

import numpy as np

var_entertainment_feat_types = ['Interests', 'Movies', 'Music']
var_personal_feat_types = ['Finance', 'Phobias']
fixed_feat_types = ['Personality', 'Health']

label_types = ['Demographic']

variables_by_type = {
    'Demographics': ['Age', 'Height', 'Weight', 'Number of siblings',
                     'Gender', 'Hand', 'Education', 'Only child', 'Home Town',
                     'Type',
                     'Home Type'],
    'Finance': ['Finances', 'Shopping centres', 'Branded clothing',
```

```

        'Entertainment spending', 'Spending on looks',
        'Spending on gadgets', 'Spending on healthy eating'],
'Health': ['Smoking', 'Alcohol', 'Healthy eating'],
'Interests': ['History', 'Psychology', 'Politics', 'Mathematics',
              'Physics', 'Internet', 'PC', 'Economy Management',
              'Biology', 'Chemistry', 'Reading', 'Geography',
              'Foreign languages', 'Medicine', 'Law', 'Cars',
              'Art exhibitions', 'Religion', 'Countryside, outdoors',
              'Dancing', 'Musical instruments', 'Writing', 'Passive sport',
              'Active sport', 'Gardening', 'Celebrities', 'Shopping',
              'Science and technology', 'Theatre', 'Fun with friends',
              'Adrenaline sports', 'Pets'],
'Movies': ['Movies', 'Horror', 'Thriller', 'Comedy', 'Romantic',
           'Sci-fi', 'War', 'Fantasy/Fairy tales', 'Animated',
           'Documentary', 'Western', 'Action'],
'Music': ['Music', 'Slow songs or fast songs', 'Dance', 'Folk',
          'Country', 'Classical music', 'Musical', 'Pop', 'Rock',
          'Metal or Hardrock', 'Punk', 'Hiphop, Rap', 'Reggae, Ska',
          'Swing, Jazz', 'Rock n roll', 'Alternative', 'Latino',
          'Techno, Trance', 'Opera'],
'Personality': ['Daily events', 'Prioritising workload',
                'Writing notes', 'Workaholism', 'Thinking ahead',
                'Final judgement', 'Reliability', 'Keeping promises',
                'Loss of interest', 'Friends versus money', 'Funniness',
                'Fake', 'Criminal damage', 'Decision making', 'Elections',
                'Self-criticism', 'Judgment calls', 'Hypochondria',
                'Empathy', 'Eating to survive', 'Giving',
                'Compassion to animals', 'Borrowed stuff',
                'Loneliness', 'Cheating in school', 'Health',
                'Changing the past', 'God', 'Dreams', 'Charity',
                'Number of friends', 'Punctuality', 'Lying', 'Waiting',
                'New environment', 'Mood swings', 'Appearance and gestures',
                'Socializing', 'Achievements', 'Responding to a serious_
↳letter',

                'Children', 'Assertiveness', 'Getting angry',
                'Knowing the right people', 'Public speaking',
                'Unpopularity', 'Life struggles', 'Happiness in life',
                'Energy levels', 'Small - big dogs', 'Personality',
                'Finding lost valuables', 'Getting up', 'Interests or_
↳hobbies',

                "Parents' advice", 'Questionnaires or polls', 'Internet_
↳usage'],
'Phobias': ['Flying', 'Storm', 'Darkness', 'Heights', 'Spiders', 'Snakes',
            'Rats', 'Ageing', 'Dangerous dogs', 'Fear of public speaking']
}

```

```
labels = variables_by_type['Demographics']
```

```

try:
    random_seed = min([StudentID1, StudentID2])
except NameError:
    random_seed = StudentID1

def which_featgroups():
    np.random.seed(random_seed)
    these_entertainments = np.random.choice(var_entertainment_feat_types, 2,
↪replace=False).tolist()
    these_personal = np.random.choice(var_personal_feat_types, 1,
↪replace=False).tolist()
    these_types = fixed_feat_types + these_personal + these_entertainments
    print('*** THESE ARE THE SELECTED TYPE OF VARIABLES:')
    for k in these_types:
        print(f'{k}')
    print('*****')
    return these_types

def which_features(these_types):
    np.random.seed(random_seed)
    these_features = []
    for type in these_types:
        if type != 'Personality':
            these_features += variables_by_type[type]
        else:
            these_features += np.random.choice(variables_by_type[type],
↪int(2 *
↪(len(variables_by_type[type]) / 3)),
↪replace=False).tolist()
    print('*** THESE ARE THE SELECTED FEATURES:')
    for ft in these_features:
        print(f'{ft}')
    print('*****')
    return these_features

these_types = which_featgroups()
these_features = which_features(these_types)

np.random.seed(random_seed)

```

```

*** THESE ARE THE SELECTED TYPE OF VARIABLES:
Personality
Health
Finance
Music
Interests

```

*** THESE ARE THE SELECTED FEATURES:

Decision making
Unpopularity
Criminal damage
Appearance and gestures
Fake
Eating to survive
Prioritising workload
Keeping promises
Dreams
Interests or hobbies
Loss of interest
Self-criticism
Socializing
Mood swings
Cheating in school
Children
Borrowed stuff
Judgment calls
Questionnaires or polls
Assertiveness
Giving
Workaholism
Writing notes
Loneliness
Achievements
Energy levels
Getting angry
God
Parents' advice
Empathy
Public speaking
Personality
Number of friends
Internet usage
Responding to a serious letter
Hypochondria
Finding lost valuables
Lying
Smoking
Alcohol
Healthy eating
Finances
Shopping centres
Branded clothing
Entertainment spending
Spending on looks

Spending on gadgets
Spending on healthy eating
Music
Slow songs or fast songs
Dance
Folk
Country
Classical music
Musical
Pop
Rock
Metal or Hardrock
Punk
Hiphop, Rap
Reggae, Ska
Swing, Jazz
Rock n roll
Alternative
Latino
Techno, Trance
Opera
History
Psychology
Politics
Mathematics
Physics
Internet
PC
Economy Management
Biology
Chemistry
Reading
Geography
Foreign languages
Medicine
Law
Cars
Art exhibitions
Religion
Countryside, outdoors
Dancing
Musical instruments
Writing
Passive sport
Active sport
Gardening
Celebrities
Shopping

Science and technology
Theatre
Fun with friends
Adrenaline sports
Pets

1.3 Importing Modules

In the following cell, import all the modules you think are necessary for doing the homework, **among the ones listed and used during the laboratories of the course**. No extra modules are allowed for reproducibility.

```
[3]: %matplotlib widget

import pandas as pd
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score, silhouette_samples
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
from matplotlib import cm
from matplotlib.lines import Line2D
from IPython.display import display # to display variables in a "nice" way
```

1.4 Exercise 1. Preparing the Dataset

In the cells below, do the following operations:

1. load the dataset “*responses_hw.csv*”;
2. create a working dataframe extracting from *responses_hw.csv* the columns corresponding to the variables in *these_features*, and randomly selecting 2/3 of the rows. Let us call this dataframe *X_df*;
3. analyze the obtained dataframe and performing cleaning/encoding operations.

```
[4]: np.random.seed(random_seed)
full_df = pd.read_csv('responses_hw.csv')
selected_columns_labels_df = full_df[these_features + labels]

selected_row = np.random.choice(np.arange(0, full_df.shape[0]), int(2 *
    ↳ (full_df.shape[0]) / 3), replace=False)
na_df = full_df.loc[selected_row].isna()

selected_columns_labels_df = selected_columns_labels_df.loc[selected_row].
    ↳ dropna()
X_df = selected_columns_labels_df.loc[:, these_features]
```

In the third point, we focused on analyzing and preparing the data to ensure its quality and suitability for further analysis. Specifically:

- **Handling Missing Values:** We removed rows with null values, which is a common practice to avoid issues during modeling and ensure the reliability of the results since in the following steps we will perform a clustering algorithm.
- **Ordinal Encoding:** For categorical variables, we applied ordinal encoding because we were confident that these variables followed a specific frequency order. This approach ensures that the encoding reflects the inherent structure of the data.

These preprocessing steps help us standardize the dataset, making it ready for further analysis or modeling. We chose to remove records having null values in fields marked as features as well as in fields marked as labels in order to be able to perform a meaningful external evaluation.

```
[5]: internet_usage_dic = {
    'no time at all': 0,
    'less than an hour a day': 1,
    'few hours a day': 2,
    'most of the day': 3
}
smoker_dic = {
    'never smoked': 0,
    'tried smoking': 1,
    'former smoker': 2,
    'current smoker': 3
}
lying_dic = {
    'never': 0,
    'only to avoid hurting someone': 1,
    'sometimes': 2,
    'everytime it suits me': 3
}
drinking_dic = {
    'never': 0,
    'social drinker': 1,
    'drink a lot': 2
}

# seleziono le feature categoriche
X_df[['Internet usage', 'Lying', 'Smoking', 'Alcohol']]

X_df['Internet usage'] = X_df['Internet usage'].apply(lambda x: internet_usage_dic[x])
X_df['Lying'] = X_df['Lying'].apply(lambda x: lying_dic[x])
X_df['Smoking'] = X_df['Smoking'].apply(lambda x: smoker_dic[x])
X_df['Alcohol'] = X_df['Alcohol'].apply(lambda x: drinking_dic[x])
X_df = X_df.astype('int16')
X_df.head()
```

```

[5]:      Decision making  Unpopularity  Criminal damage  Appearance and gestures  \
183          3          4          1          5
59           1          5          1          4
271          4          3          5          3
341          5          5          5          5
520          2          3          2          5

      Fake  Eating to survive  Prioritising workload  Keeping promises  Dreams  \
183     1          2          4          4          4
59      5          2          4          4          4
271     2          2          3          4          3
341     3          4          1          2          3
520     2          3          3          3          4

      Interests or hobbies  ...  Passive sport  Active sport  Gardening  \
183          4  ...          3          5          3
59          5  ...          1          1          1
271          4  ...          4          5          1
341          5  ...          5          5          3
520          4  ...          4          4          3

      Celebrities  Shopping  Science and technology  Theatre  Fun with friends  \
183          1          3          5          4          5
59          1          1          1          1          3
271          4          4          3          3          5
341          4          5          5          1          5
520          3          5          3          3          5

      Adrenaline sports  Pets
183          5          5
59          1          1
271          5          5
341          5          2
520          3          3

[5 rows x 99 columns]

```

1.5 Exercise 2. Analyzing the Variance and the PCs

In the cells below, do the following operations:

1. create two new dataframes from `X_df` applying a `StandardScaler` and a `MinMaxScaler`. Call these new dataframes as `Xstd_df` and `Xmm_df`, respectively;
2. compute the variance of all the features in `X_df`, `Xstd_df`, and `Xmm_df` and **comment the results**;
3. compute all the n Principal Components (PCs) for each dataset `X_df`, `Xstd_df`, and `Xmm_df`. Then, visualize the curves of the cumulative explained variances and **comment the results**.


```
[6]: scaler = StandardScaler()
minmax = MinMaxScaler()


Xstd_df = pd.DataFrame(scaler.fit_transform(X_df), columns=X_df.columns)
Xmm_df = pd.DataFrame(minmax.fit_transform(X_df), columns=X_df.columns)

fig, axs = plt.subplots(1, 3, figsize = (10, 4))

dfs = [X_df, Xstd_df, Xmm_df]
titles = ['Original data', 'Standardized data', 'MinMax scaled data']

for title, df, ax in zip(titles, dfs, axs):
    var_s = df.var().sort_values()
    ax.barh(var_s.index, var_s.values)
    ax.set_yticks([])
    ax.set_title(title)
    ax.set_xlabel("Variance")
    ax.set_ylabel("Features")

plt.tight_layout()
plt.show()
```



Variance of features in pre-processed data depends on the type of pre-processing they are subject to:

- features of *standardized data*, as expected, have unit variance due to the standardization process they've gone through;
- features of *MinMax scaled data* have a variance whose shape resembles the shape of variance of the features of the original data; however being features rescaled in the range (0, 1) their variance is scaled by a factor that depends on the square of the range the feature covered

before being standardized:

$$\sigma^2 = \sum_{i=1}^n (x_i - \bar{x})^2$$

$$\sigma_{mm}^2 = \sum_{i=1}^n \left(\frac{x_i - x_{min} - \bar{x} + x_{min}}{x_{max} - x_{min}} \right)^2 = \frac{\sigma^2}{R^2}$$

where $R = x_{max} - x_{min}$ is the length of the range covered by the feature. Since most features cover the same range (1 to 5) their variance decreases by similar factors, hence the resulting shape is similar.

In the following code cell we show that 94 features over 99 have values ranging from 1 to 5, hence range R is 4 for most features.

```
[7]: pd.DataFrame([X_df.min(), X_df.max()]).T.apply(lambda x: (x[0], x[1]), axis=1).
      ↪ value_counts()
```

```
[7]: (1, 5)    94
      (0, 3)     2
      (1, 3)     1
      (0, 2)     1
      (2, 5)     1
      Name: count, dtype: int64
```

```
[8]: pca_x_df = PCA()
      pca_xmm_df = PCA()
      pca_xstd_df = PCA()

      pca_x_df.fit(X_df.values)
      pca_xmm_df.fit(Xmm_df.values)
      pca_xstd_df.fit(Xstd_df.values)

      fig, ax = plt.subplots(1, 1, figsize=(8, 5))
      pca_objs = [pca_x_df, pca_xstd_df, pca_xmm_df]

      for pca_obj, title in zip(pca_objs, titles):
          ax.plot(np.cumsum(pca_obj.explained_variance_ratio_), label=title)


      ax.set_xlabel('Principal components')
      ax.set_ylabel('Cumulative explained variance fraction')
      ax.set_ylim(0, 1)
      ax.set_xlim(0, pca_obj.n_features_in_)
      # Set ticks and labels
      ax.set_xticks(np.arange(pca_obj.n_features_in_))
      PCs_labels=[f'PC{i + 1}' for i in range(pca_obj.n_features_in_)]
      [l.set_visible(False) for (i,l) in enumerate(ax.xaxis.get_ticklabels()) if i % 5
      ↪ != 0]
      ax.set_xticklabels(PCs_labels, rotation=45, ha='center', label = title)
```

```

ax.grid()
ax.legend()
# Draw a line corresponding to the threshold of the following exercise
ax.axhline(0.33, c='r')

fig.tight_layout()

```



In the previous figure, one can notice that the curves are similar for original data and MinMax scaled data: this is due to the considerations that were previously made for the distribution of variances among the features of data undergoing different preprocessing steps.

- For *Standardized data*, variance is more evenly distributed across features.
- For *MinMax scaled data*, being the variance of features rescaled by a factor that is the same for pretty much all the features, the same features as in the original data explain most of the variation in the data.

1.6 Exercise 3. Dimensionality Reduction and PC Interpretation

In the cells below, do the following operations:

1. For each one of the two dataframes *Xstd_df*, and *Xmm_df*, compute a new PCA for performing a dimensionality reduction with respect to m dimensions. The value of m must be

$$m = \min\{m', 5\},$$

where m' is the value required for obtaining 33% of the total variance.

2. For both the cases, visualize all the PCs and give a name/interpretation to them. **Comment and motivate your interpretations.** If possible, **compare the differences among the results obtained** for *Xstd_df* and *Xmm_df*.
3. Perform the score graph for both the cases (*std* and *mm*). If $m > 3$, plot the score graph with respect to the first 3 PCs. All the **plots must show the names of the PCs on the axes** for better understanding the results.
4. **Optional:** plot more score graphs, coloring the dots with respect to any label in the list *labels* that you believe can be interesting. **Comment and analyze this optional plots.**

```
[9]: threshold = 0.33

m_xmm_threshold = (np.cumsum(pca_xmm_df.explained_variance_ratio_) < threshold).
    ↳sum() + 1
m_std_threshold = (np.cumsum(pca_xstd_df.explained_variance_ratio_) <
    ↳threshold).sum() + 1


m_xmm = min(m_xmm_threshold, 5)
m_std = min(m_std_threshold, 5)

pca_xmm_df = PCA(n_components=m_xmm)
pca_xstd_df = PCA(n_components=m_std)
pca_xmm_df.fit(Xmm_df.values)
pca_xstd_df.fit(Xstd_df.values)

m_xmm, m_std
```

```
[9]: (5, 5)
```

```
[10]: fig, axs = plt.subplots(1, 2, figsize=(8, 4), sharey=True)
for pca_obj, m, title, ax in zip([pca_xstd_df, pca_xmm_df, ], [m_std, m_xmm],
    ↳titles[1:], axs):
    PCs_labels = [f'PC{i + 1}' for i in range(m)]
    x = np.arange(m)
    ax.bar(x, height=pca_obj.explained_variance_ratio_[0:m])
    ax.set_xticks(x)
    ax.set_xticklabels(PCs_labels)
    ax.set_title(title)
    ax.set_xlabel("Principal component")
    ax.set_ylabel("Explained variance fraction")
plt.tight_layout()
```



```
[11]: def get_question(feature_name):
        columns = pd.read_csv('columns_hw.csv', header=None)
        columns['question'] = columns[0].apply(lambda x: x.split(',')[0])
        columns['name'] = columns[0].apply(lambda x: x.split(',')[1])
        columns = columns.drop(columns=0)
        return columns.loc[columns['name'] == feature_name]['question']
```

To better understand the PCA analysis, we color the bars of the attributes according to the category of the feature. This visualization helps us identify patterns and relationships between features and categories more clearly, enhancing the interpretability of the PCA results.

```
[12]: set2 = cm.Set2.colors

type_colors = {
    'Personality': set2[0],
    'Health': set2[1],
    'Finance': set2[2],
    'Interests': set2[3],
    'Music': set2[4]
}

these_types_col = []
for t in these_types:
    if t != 'Personality':
        these_types_col += [t]*len(variables_by_type[t])
    else:
        these_types_col += [t]*int(2 * (len(variables_by_type[t]) / 3))

types_df = pd.DataFrame()
```

```
types_df['type'] = these_types_col
types_df['feature'] = these_features
types_df['color'] = types_df['type'].apply(lambda x: type_colors[x])
```

```
[13]: def plot_feature_contribution_pc(X, pca, ii, ax, title, eps_coefficient=1,
    verbose=False):

    eps = eps_coefficient * np.sqrt(1 / pca.n_features_in_)
    color_types = types_df['color']
    type_colors_legend = [Line2D([0], [0], color=type_colors[k])
        for k in type_colors.keys()]

    ax.barh(np.arange(pca.n_features_in_),
        pca.components_[ii, :], color=color_types)
    # --- RED LINE DENOTING THE THRESHOLD [-eps, +eps] -----
    ax.plot([eps, eps], [-0.5, pca.n_features_in_ - 0.5], 'red')
    ax.plot([-eps, -eps], [-0.5, pca.n_features_in_ - 0.5], 'red')
    # -----
    ax.set_yticks(ticks=np.arange(pca.n_features_in_),
        labels=X.columns.to_list())
    ax.set_xlim(-0.275, 0.275)
    ax.set_ylim(ax.get_yticks()[0] - 0.5, ax.get_yticks()[-1] + 0.5)
    ax.set_title(f'PC{ii+1} - {title}')
    ax.grid()
    ax.legend(type_colors_legend, [k for k in type_colors.keys()], loc="lower_
    right")


    # THE SELECTION OF THE FEATURES WITH CONTRIBUTE GREATER THAN THE THRESHOLD
    ind_great_pos_PCii = np.argwhere(
        pca.components_[ii, :] >= eps).flatten()
    ind_great_neg_PCii = np.argwhere(
        pca.components_[ii, :] <= -eps).flatten()


    for i in ind_great_pos_PCii:
        ax.get_yticklabels()[i].set_color('r')
    for i in ind_great_neg_PCii:
        ax.get_yticklabels()[i].set_color('b')


    great_pos_PCii = [X.columns[i] for i in ind_great_pos_PCii]
    great_neg_PCii = [X.columns[i] for i in ind_great_neg_PCii]
    if verbose:
        print('')
        print(f'***** PC{ii+1} *****')
        print(f'HIGH-VALUED POSITIVE COMPONENTS: {great_pos_PCii}')
        print('')
        print(f'HIGH-VALUED NEGATIVE COMPONENTS: {great_neg_PCii}')
```


```
print('*****')
print('')
```


```
[14]: for i in range(5):
        fig, axs = plt.subplots(1, 2, figsize=(10, 13), sharex=True)
        plot_feature_contribution_pc(Xmm_df, pca_xmm_df, i, axs[0], "MinMax Scaled_↵
        ↵Data")
        plot_feature_contribution_pc(Xstd_df, pca_xstd_df, i, axs[1], "Standardized_↵
        ↵Data")
        fig.tight_layout()
```











The discussion on the interpretation of the PCs is postponed to the following sections. We now compare the contribution of the features to the principal components that arise from data being scaled according to MinMax or standardized:

1. (PC1) The contribution of the features are similar.
2. (PC2) The shapes are similar, but standardized data's feature that positively contribute have a lower importance.
3. (PC3) The contribution of the features are similar.
4. (PC4) The contribution of the features are similar, but one is the opposite of the other.
5. (PC5) This is the only PC where the contribution of the features are different.

1.6.1 MinMax Scaled Data

For MinMax scaled data the following interpretations are given for the first 5 principal components.

1. *Transgression vs. Culture*:
 - (+) Interest in many areas of culture, interest in music genres that are commonly associated with culture, healthy behaviours, personality traits that suggest diligence;
 - (-) Personality traits that suggest criminal behaviour;
2. *Pop vs. Alternative*:
 - (+) Interest in alternative music, interest in STEM;
 - (-) Interest in social activities, carefree spending patterns, interest in commercial music, personality traits of an extrovert;
3. *Introvert vs. Tech Enthusiast*:
 - (+) Interest in STEM, carefree spending patterns, focused on interest and hobbies;
 - (-) Personality traits of an introvert;
4. *Analytical vs. Creative*:
 - (+) Artsy interests, bad health behaviours;
 - (-) Interest in STEM, interest in finance, prioritizes workload;
5. *Social Sciences vs. Life Sciences*:
 - (+) Interest in life sciences;
 - (-) Interest in social sciences.


```
[15]: pc_names_xmm = ['Transgression vs. Culture',
                    'Pop vs. Alternative',
                    'Introvert vs. Tech Enthusiast',
                    'Quantitative Reasoning vs. Humanist Culture',
                    'Social Sciences vs. Life Sciences']
```

```
[16]: Y_mm = pca_xmm_df.transform(Xmm_df.values)

# MAKE THE 3D SCORE GRAPH
sg_3d = plt.figure(figsize=(14,6))
ax_sg_3d = sg_3d.add_subplot(111, projection='3d')
ax_sg_3d.scatter(Y_mm[:, 0], Y_mm[:, 1], Y_mm[:, 2], s=2)
plt.title('Survey - SCORE GRAPH')

ax_sg_3d.set_xlabel(pc_names_xmm[0])
ax_sg_3d.set_ylabel(pc_names_xmm[1])
ax_sg_3d.set_zlabel(pc_names_xmm[2])
plt.grid()
```

```
plt.show()
plt.tight_layout()
```



```
[17]: df_label = selected_columns_labels_df.loc[:, labels]

def plot_score_graph_2D(field, Y, pc_names, alpha=1):
    types = df_label[field].unique()
    cmap = plt.get_cmap('viridis')
    colors = cmap(np.linspace(0, 1, types.size))
    color_dict = {type: color for type, color in zip(types, colors)}
    color_series = df_label[field].apply(lambda x: color_dict[x])
    colors_legend = [Line2D([0], [0], color=k) for k in color_dict.values()]

    # MAKE THE 3D SCORE GRAPH
    fig, ax = plt.subplots()
    ax.scatter(Y[:, 0], Y[:, 1], s=20, c=color_series.values, alpha=alpha)

    ax.set_xlabel(pc_names[0])
    ax.set_ylabel(pc_names[1])

    ax.grid()
    ax.legend(colors_legend, [k for k in color_dict.keys()])
    plt.tight_layout()
    plt.show()
    return ax

def plot_score_graph(field, Y, pc_names, alpha=1):
    types = df_label[field].unique()
    cmap = plt.get_cmap('viridis')
    colors = cmap(np.linspace(0, 1, types.size))
```

```

color_dict = {type: color for type, color in zip(types, colors)}
color_series = df_label[field].apply(lambda x: color_dict[x])
colors_legend = [Line2D([0], [0], color=k) for k in color_dict.values()]

# MAKE THE 3D SCORE GRAPH
sg_3d = plt.figure(figsize=(10,6))
ax_sg_3d = sg_3d.add_subplot(111, projection='3d')
ax_sg_3d.scatter(Y[:, 0], Y[:, 1], Y[:, 2], s=10, c=color_series.values,
↪alpha=alpha)

ax_sg_3d.set_xlabel(pc_names[0])
ax_sg_3d.set_ylabel(pc_names[1])
ax_sg_3d.set_zlabel(pc_names[2])

plt.grid()
plt.legend(colors_legend, [k for k in color_dict.keys()])
plt.title(field)
plt.tight_layout()
plt.show()
return ax_sg_3d

```

In order to plot score graphs where points are colored according to a certain label, since many labels have lots of different values and we want to avoid using too many colours which may make the figures difficult to read, we split continuous features or features which have too many categorical values into ranges.

```


[18]: age = df_label['Age']

age = age.value_counts()

fig, ax = plt.subplots()

ax.axhline(0.5, c='red')
ax.plot(np.sort(df_label['Age'].unique()), age.sort_index().values.cumsum()/age.
↪sort_index().values.sum(), drawstyle='steps')
ax.grid(visible=True)
ax.set_ylim(0, 1)
ax.set_xlim(df_label['Age'].min(), df_label['Age'].max())
ax.set_xlabel('Age')
ax.set_ylabel('Fraction')
ax.set_title('Empirical cumulative distribution function of the age')
plt.tight_layout()
plt.show()

```



Having plotted the empirical cumulative distribution of the ages, we chose to split ages in two ranges: age less or equal than 20 and age greater than 20 since 20 is the median age.

```
[19]: df_label['AgeGt20'] = df_label.Age.apply(lambda x: "<= 20" if x <= 20 else ">20")
```

Since height and weight did not show any particular pattern with respect to the 3 first principal components, we chose to compute the BMI and to assign labels to distinct ranges of the BMI.

```
[20]: def get_bmi_label(bmi):
    if bmi <= 18.5:
        return "Thinness"
    elif bmi <= 25:
        return "Ideal Weight"
    elif bmi <= 30:
        return "Overweight"
    return "Obese"

arr = np.array(df_label.values[:, :3], dtype=np.float32)
df_label['BMI'] = arr[:, 2]/np.square(arr[:, 1]/100)
```



```
df_label['BMI Label'] = df_label['BMI'].apply(get_bmi_label)
```


Since education had lots of distinct labels that were not so meaningful, we chose to group them into the following ranges:



- *undergrad*, for people who did not continue studies past college;
- *degree*, for people who continued studies past college.


```
[21]: def get_education_label(education):
        if education in ['college/bachelor degree', 'masters degree', 'doctorate_
        degree']:
            return 'degree'
        return "undergrad"

df_label['Education label'] = df_label['Education'].apply(get_education_label)
```

```
[22]: plot_score_graph('AgeGt20', Y_mm, pc_names_xmm)
        plot_score_graph('Gender', Y_mm, pc_names_xmm)
        plot_score_graph("BMI Label", Y_mm, pc_names_xmm)
        plot_score_graph('Education label', Y_mm, pc_names_xmm)
```







[22]: `<Axes3D: title={'center': 'Education label'}, xlabel='Transgression vs. Culture', ylabel='Pop vs. Alternative', zlabel='Introvert vs. Tech Enthusiast'>`

Unlike it was expected, no labels but gender seem to have a distribution that is highly correlated to the values of the first 3 principal components. One can notice that men tend to be transgressive, alternative and keen on technology, while women engage in cultural activities, pop culture and have traits associated with introversion.

1.6.2 Standardized Data

For standardized data the following interpretations are given for the first 5 principal components.

1. *Transgression vs. Culture*:
 - (+) Interest in many areas of culture, interest in music genres that are commonly associated with culture, healthy behaviours, personality traits that suggest diligence;
 - (-) Personality traits that suggest criminal behaviour;
2. *Extrovert vs. Introvert*:
 - (+) Feelings of loneliness and self-criticism, interest in Metal music;
 - (-) Interest in social activities, carefree spending patterns, interest in commercial music, personality traits of an extrovert;
3. *Pop vs. Alternative*:
 - (+) Interest in alternative music, interest in STEM;
 - (-) Interest in social activities, carefree spending patterns, interest in commercial music;
4. *Creative vs. Analytical*:
 - (+) Interest in STEM, interest in finance, prioritizes workload;
 - (-) Artsy interests, bad health behaviours;
5. *Modern Recluse vs. Social Butterfly*:


- (+) Personality traits of an extrovert;
- (-) Geeky interests, loneliness, high usage of internet, carefree spending habits.

```
[23]: pc_names_std = [
    'Transgression vs. Culture',
    'Extrovert vs. Introvert',
    'Pop vs. Alternative',
    'Analytical vs. Creative',
    'Modern recluse vs. Social butterfly'
]
```



```
[24]: Y_std = pca_xstd_df.transform(Xstd_df.values)



# MAKE THE 3D SCORE GRAPH
sg_3d = plt.figure(figsize=(8, 8))
ax_sg_3d = sg_3d.add_subplot(111, projection='3d')
ax_sg_3d.scatter(Y_std[:, 0], Y_std[:, 1], Y_std[:, 2], s=10)
plt.title('Survey - SCORE GRAPH')
ax_sg_3d.set_xlabel(pc_names_std[0])
ax_sg_3d.set_ylabel(pc_names_std[1])
ax_sg_3d.set_zlabel(pc_names_std[2])
plt.grid()
plt.show()
```

Survey - SCORE GRAPH



```
[25]: plot_score_graph('AgeGt20', Y_std, pc_names_std)
      plot_score_graph('Gender', Y_std, pc_names_std)
      plot_score_graph("BMI Label", Y_std, pc_names_std)
      plot_score_graph('Education label', Y_std, pc_names_std)
```





[25]: `<Axes3D: title={'center': 'Education label'}, xlabel='Transgression vs. Culture', ylabel='Extrovert vs. Introvert', zlabel='Pop vs. Alternative'>`

Unlike it was expected, no labels but gender seem to have a distrubution that is highly correlated

to the values of the first 3 principal components. One can notice that men tend to be transgressive, alternative and extrovert, while women engage in cultural activities, pop culture and have traits associated with introversion.

1.7 Exercise 4. k -Means

In the cells below, do the following operations:

1. For each one of the two datasets (*std* and *mm*), run the k -Means for clustering the data. In particular, **use the silhouette score for identify the best value for $k \in \{3, \dots, 10\}$.**
2. Plot the score graphs of exercise 3.3, adding the centroids of the cluster.
3. Observing the centroids coordinates in the PC space, **give a name/interpretation to them**, exploiting the names you assigned to the PCs. **Comment and motivate your interpretations.**

```
[26]: random_state = random_seed

def best_k(Y, verbose=False):
    # INITIALIZE SOME LISTS TO STORE THE TEMPORARY RESULTS AND, THEN, MAKE
    # COMPARISONS
    km_list = []
    silcoeff_list = []
    k_list = list(range(3, 11))

    # START THE FOR-CYCLE TO RUN THE k-MEANS AND MEASURING THE SILHOUETTE
    # COEFFICIENT
    for i in range(len(k_list)):
        if verbose:
            print(f'***** START k-MEANS WITH k={k_list[i]}')
            print('Computing...')
        km_list.append(KMeans(k_list[i], random_state=min([StudentID1,
            StudentID2])))
        km = km_list[i]
        # attenzione il fit va fatto sui dati trasformati con il pca

        km.fit(Y)
        silcoeff_list.append(silhouette_score(Y, km.labels_))
        if verbose:
            print(f'***** END k-MEANS WITH k={k_list[i]}')
            print('')

    # FIND THE BEST VALUE OF k AND THE BEST KMeans OBJECT
    i_best = np.argmax(silcoeff_list)
    k = k_list[i_best]
    km = km_list[i_best]
```



```

# VISUALIZE THE RESULT
if verbose:
    print('')
    print('')
    print('***** RESULTS OF THE SEARCH... *****')
    print(f'BEST SILHOUETTE SCORE: {np.max(silcoeff_list)} --> k = {k}')
    print('*****')
return k, km, k_list, silcoeff_list

def plot_sil_score(k_list, silcoeff_list):
    plt.figure()
    plt.plot(k_list, silcoeff_list)
    plt.grid(visible=True)
    plt.xlabel('k')
    plt.ylabel('Silhouette score')
    plt.title('Silhouette Scores for Different Values of k')
    plt.xlim(min(k_list), max(k_list))
    plt.show()
    return

def plot_score_graph_cluster(field, Y, pc_names, k, km):
    ax_sg_3d = plot_score_graph(field, Y, pc_names, 0.3)

    for kk in range(k):
        ax_sg_3d.text(km.cluster_centers_[kk, 0], km.cluster_centers_[kk, 1],
            ↪km.cluster_centers_[kk, 2], f'clust.{kk + 1}')
        ax_sg_3d.scatter(km.cluster_centers_[ :, 0], km.cluster_centers_[ :, 1], km.
            ↪cluster_centers_[ :, 2], color='black', alpha=1)
    return

def plot_score_graph_cluster_2D(field, Y, pc_names, k, km):

    ax = plot_score_graph_2D(field, Y, pc_names, 0.3)

    for kk in range(k):
        ax.text(km.cluster_centers_[kk, 0], km.cluster_centers_[kk, 1],
            ↪f'clust.{kk + 1}')
        ax.scatter(km.cluster_centers_[ :, 0], km.cluster_centers_[ :, 1],
            ↪color='black', alpha=1)
    return

def plot_centroids_bar(k, km, Y, pc_names, title):
    maxs_y = Y.max(axis=0) # <-- TODO!!
    mins_y = Y.min(axis=0)

    fig_centroids, ax_centroids = plt.subplots(1, 3, figsize=(12, 5))

```

```

    for ii in range(k):
        ax_centroids[ii].bar(np.arange(km.cluster_centers_.shape[1]), maxs_y,
        color='blue', alpha=0.15)
        ax_centroids[ii].bar(np.arange(km.cluster_centers_.shape[1]), mins_y,
        color='blue', alpha=0.15)
        ax_centroids[ii].bar(np.arange(km.cluster_centers_.shape[1]), km.
        cluster_centers_[ii, :])
        ax_centroids[ii].set_xticks(ticks=np.arange(km.cluster_centers_.
        shape[1]))
        ax_centroids[ii].set_xticklabels(labels=pc_names, rotation=45,
        ha='right')
        ax_centroids[ii].grid(visible=True, which='both')
        ax_centroids[ii].set_title(f'{title} - CENTROID {ii+1}')

plt.tight_layout()
plt.show()
return


```

1.7.1 minmax

```

[27]: k_mm, km_mm, k_list_mm, silcoeff_list_mm = best_k(Y_mm)
      plot_sil_score(k_list_mm, silcoeff_list_mm)


```




```
[28]: df_label['cluster_mm'] = km_mm.labels_ + 1
      plot_score_graph_cluster('cluster_mm', Y_mm, pc_names_xmm, k_mm, km_mm)
      plot_centroids_bar(k_mm, km_mm, Y_mm, pc_names_xmm, 'MinMax Scaled')
```


/tmp/ipykernel_70/3834587576.py:33: RuntimeWarning: More than 20 figures have been opened. Figures created through the pyplot interface (`matplotlib.pyplot.figure`) are retained until explicitly closed and may consume too much memory. (To control this warning, see the rcParam `figure.max_open_warning`). Consider using `matplotlib.pyplot.close()`.



```
sg_3d = plt.figure(figsize=(10,6))
```






```
[29]: plot_score_graph_cluster('Gender', Y_mm, pc_names_xmm, k_mm, km_mm)
      plot_score_graph_cluster('Education label', Y_mm, pc_names_xmm, k_mm, km_mm)
      plot_score_graph_cluster('AgeGt20', Y_mm, pc_names_xmm, k_mm, km_mm)
      plot_score_graph_cluster('BMI Label', Y_mm, pc_names_xmm, k_mm, km_mm)
```







We repeated the plot of the score graphs with the previously considered labels encoded into the point color. There is no association between cluster centers and label of the points. Cluster 2 seems to be associated to males, while cluster 3 to females. Cluster 1 is somewhere in between genders.


```
[30]: cluster_names_mm = ['Cultured',
                          'Popular Transgressive',
                          'Introvert Alternative'
                          ]
```

1. Cluster 1 has been named *Cultured* since it has a significant component only in the positive direction of PC1.
2. Cluster 2 has been named *Popular Transgressive* since it has a significant components in the negative directions of PC1 and PC2.
3. Cluster 3 has been named *Introvert Alternative* since it has significant components in the negative direction of PC3 and in the positive direction of PC2.

We did not consider the last two PCs while assigning names to the clusters since their values are negligible and close to zero for all clusters.

1.7.2 Std


```
[31]: k_std, km_std, k_list_std, silcoeff_list_std = best_k(Y_std)
      plot_sil_score(k_list_std, silcoeff_list_std)
```





```
[32]: km_std.cluster_centers_#[:, :3]
```



```
[32]: array([[ 2.88340634,  0.53487304,  0.68475262, -0.14030705, -0.0359591 ],  
          [-2.56529436,  1.15519797,  1.28484775,  0.38813167, -0.07724426],  
          [-0.53822748, -1.42333101, -1.66470884, -0.18746215,  0.09534961]])
```

```
[33]: df_label['cluster_std'] = km_std.labels_ + 1  
      plot_score_graph_cluster('cluster_std', Y_std, pc_names_std, k_std, km_std)  
      plot_centroids_bar(k_std, km_std, Y_std, pc_names_std, 'Std Scaled')
```



```
[34]: plot_score_graph_cluster('Gender', Y_std, pc_names_std, k_std, km_std)
      plot_score_graph_cluster('Education label', Y_std, pc_names_std, k_std, km_std)
      plot_score_graph_cluster('AgeGt20', Y_std, pc_names_std, k_std, km_std)
      plot_score_graph_cluster('BMI Label', Y_std, pc_names_std, k_std, km_std)
```



We repeated the plot of the score graphs with the previously considered labels encoded into the point color. There is no association between cluster centers and label of the points. Cluster 2 seems to be associated to males, while cluster 3 to females. Cluster 1 is somewhere in between genders.

```
[35]: cluster_names_std = ['Cultured',
                           'Transgressive Introvert Alternative',
                           'Popular Extrovert'
                           ]
```

1. Cluster 1 has been named *Cultured* since it has a significant component only in the positive direction of PC1.
2. Cluster 2 has been named *Transgressive Introvert Alternative* since it has a significant components in the negative directions of PC1, positive direction of PC2 and positive direction of PC3.
3. Cluster 3 has been named *Popular Extrovert* since it has significant components in the negative directions of PC2 and PC3.

We did not consider the last two PCs while assigning names to the clusters since their values are negligible and close to zero for all clusters.

1.8 Exercise 5. Cluster Evaluations

In the cells below, do the following operations:

1. For each one of the two datasets (*std* and *mm*), perform an **external evaluation** of the clustering obtained at exercise 4.1 with respect to one or more labels in the list *labels*. **Comment the results, comparing the evaluation with the interpretation you gave at exercise 4.3.**
2. For each one of the two datasets (*std* and *mm*), perform an **internal evaluation** of each cluster, with respect to the silhouette score. **Comment the results.**

```
[36]: def external_evaluation(Y, k, km, field, cluster_names, title="External_
      ↪evaluation"):

    df_cluster = pd.DataFrame(km.predict(Y), columns=['cluster'], dtype=np.
    ↪int16)
    df_cluster[field] = df_label[field]

    labels = df_label[field].unique()

    fig, ax = plt.subplots(1, k, figsize=(4*k, 4), sharey=True)

    counts_total = dict(df_cluster[field].value_counts())

    for i in range(k):
        df_cluster_i = dict(df_cluster.loc[df_cluster['cluster'] == i][field].
    ↪value_counts())
        ax[i].grid(visible = True)
        ax[i].bar(labels, [df_cluster_i.get(l, 0) for l in labels])
        ax[i].set_title(f'Cluster {i+1} - {cluster_names[i]}')
        ax[i].set_xticks(ax[i].get_xticks())
        ax[i].set_xticklabels(ax[i].get_xticklabels(), rotation=45, ha='right')
```

```

ax[i].set_ylabel("# Points")

plt.suptitle(title)
plt.tight_layout()
plt.show()
return

def internal_evaluation(cluster_names, Y, k, km):
    sil_scores = silhouette_samples(Y, km.labels_)
    cluster_sil = [sil_scores[km.labels_ == i].mean() for i in range(k)]

    return pd.DataFrame(cluster_sil, index=cluster_names,
        columns=['sil_score'])#.sort_values('sil_score', ascending=False)


```

1.8.1 External evaluation


```

[37]: for label in ['Gender', 'Education', 'BMI Label']:
        external_evaluation(Y_mm, k_mm, km_mm, label, cluster_names_mm, f'External_
        evaluation Minmax - {label}')
        external_evaluation(Y_std, k_std, km_std, label, cluster_names_std,
        f'External evaluation Std - {label}')


```




External evaluation Minmax - Education




External evaluation Std - Education



External evaluation Minmax - BMI Label





From the previous barplots one can notice that distribution of labels is similar across different clusters, even though we may have expected a different behaviours for some of the chosen labels:

1. From the graphs it seemed that in the standardized data case males and females were to belong to different clusters, but it's not true (this may be due to the fact that there are 1.5 times more females than males in the dataset).
2. Even though in the minmax scaled data case as well as in the standardized data case there is a cluster of supposedly cultured people, they don't exhibit an average higher level of education when compared to the other clusters.

1.8.2 Internal Evaluation

As one could expect from the average silhouette score of the points, also the average silhouette score by cluster is not a large value and as seen in the plot of the clusters by label the cluster are not so distinguished from one another. This holds true for MinMax scaled data as well as for standardized data, but in MinMax scaled data the difference is a little more evident between clusters.

1.8.3 MinMax scaled data

```
[38]: internal_evaluation(cluster_names_mm, Y_mm, k_mm, km_mm)
```

```
[38]:
           sil_score
Cultured          0.161822
Popular Transgressive  0.202534
Introvert Alternative  0.122716
```

1.8.4 Standardized data

```
[39]: internal_evaluation(cluster_names_std, Y_std, k_std, km_std)
```

```
[39]:
           sil_score
Cultured          0.141857
Transgressive Introvert Alternative  0.146980
Popular Extrovert  0.200965
```